

PROBLEM 1: CONVEX FUNCTIONS

Answer Q1A: Let us derive this from the probability of one data point, and assume IID data:

$$\begin{aligned}
 P(Y = y|X = \mathbf{x}) &= \sigma(\theta^T \mathbf{x})^y \cdot [1 - \sigma(\theta^T \mathbf{x})]^{(1-y)} \\
 \implies L(\theta|y) &= \prod_{i=1}^n P(Y = y_i|X = \theta_i) \\
 &= \prod_{i=1}^n p(\mathbf{x}_i)^{y_i} \cdot \left(1 - p(\mathbf{x}_i)\right)^{1-y_i}
 \end{aligned} \tag{1}$$

Now take logs and substitute $p(\mathbf{x})$ with the exponent form to give:

$$\begin{aligned}
 \log L(\theta|y) &= \sum_{i=1}^n y_i \log \left(\frac{1}{1 + e^{\theta^T \mathbf{x}_i}} \right) + (1 - y_i) \log \left(\frac{e^{-\theta^T \mathbf{x}_i}}{1 + e^{-\theta^T \mathbf{x}_i}} \right) \\
 \implies &= \sum_{i=1}^n y_i \left[\log \left(\frac{1}{1 + e^{\theta^T \mathbf{x}_i}} \right) - \log \left(\frac{e^{-\theta^T \mathbf{x}_i}}{1 + e^{-\theta^T \mathbf{x}_i}} \right) \right] + \log \left(\frac{e^{-\theta^T \mathbf{x}_i}}{1 + e^{-\theta^T \mathbf{x}_i}} \right) \\
 \implies &= \sum_{i=1}^n y_i \theta_i^T \mathbf{x}_i - \log(1 + e^{\theta^T \mathbf{x}_i}) \\
 \implies -\log L(\theta|y) &= \sum_{i=1}^n -y_i \theta^T \mathbf{x}_i + \log(1 + e^{\theta^T \mathbf{x}_i}), \text{ as required. } \quad \square
 \end{aligned} \tag{2}$$

Answer Q1B: First, let:

$$p_i = p(\mathbf{x}_i) = \sigma(\theta^T \mathbf{x}_i), \quad a_i = y_i \log(p_i) + (1 - y_i) \log(1 - p_i), \quad t_i = \theta^T \mathbf{x}_i$$

Now we have:

$$\begin{aligned}
 L(\theta) &= - \sum_{i=1}^n \left(y_i \log p_i + (1 - y_i) \log(1 - p_i) \right) \\
 \implies \nabla_{\theta} a_i &= \frac{y_i}{p_i} (p_i(1 - p_i)) \nabla_{\theta} t_i - \frac{(1 - y_i)}{(1 - p_i)} (p_i(1 - p_i)) \nabla_{\theta} t_i \\
 \implies \nabla_{\theta} a_i &= y_i(1 - p_i) \mathbf{x}_i - (1 - y_i) p_i \mathbf{x}_i = (y_i - p_i) \mathbf{x}_i.
 \end{aligned} \tag{3}$$

Finally we have:

$$\begin{aligned}
 \nabla_{\theta} \log L(\theta|y) &= \sum_{i=1}^n (p_i - y_i) \mathbf{x}_i \equiv \mathbf{X}^T (\mathbf{p} - \mathbf{y}) \\
 \implies -\nabla_{\theta} \log L(\theta|y) &= -\mathbf{X}^T (\mathbf{y} - \sigma(\theta)), \text{ as required. } \quad \square
 \end{aligned} \tag{4}$$

Answer Q1C: From Equation 3 we derived the first gradient which can now be utilized to derive the Hessian of the negative log-likelihood. Thus we can now give the the elements of the Hessian matrix as follows for elements k, j (row index, column index):

$$\begin{aligned} [\nabla_{\theta}^2 a_i]_{kj} &= \frac{d}{d\theta_i} \left[(p_i - y_i) x_i(j) \right] \\ &= (p_i(1 - p_i)) \cdot x_i(k) \cdot x_i(j) \\ &\equiv [\mathbf{x}_i \mathbf{d}_i \mathbf{x}_i^T]_{kj} \\ \implies \nabla_{\theta}^2 a_i &= \mathbf{x}_i \mathbf{d}_i \mathbf{x}_i^T \end{aligned} \quad (5)$$

Using the same notation for the diagonal matrix shown in the question this implies:

$$-\nabla_{\theta}^2 \log L(\theta|y) = \sum_{i=1}^n \mathbf{x}_i \mathbf{d}_i \mathbf{x}_i^T = \mathbf{X}^T \mathbf{D}(\theta) \mathbf{X}, \text{ as required. } \quad \square \quad (6)$$

Answer Q1D: Using the Equation 6 and the positive definite of our Diagonal matrix $\mathbf{D}(\theta)$ we can simply derive the following:

$$\begin{aligned} -\nabla_{\theta}^2 \log L(\theta|y) &= \mathbf{X}^T \mathbf{D}(\theta) \mathbf{X} \\ \implies \mathbf{x}^T \mathbf{X}^T \mathbf{D}(\theta) \mathbf{X} \mathbf{x} &= \mathbf{y}^T \mathbf{D}(\theta) \mathbf{y} > 0. \end{aligned} \quad (7)$$

where $\mathbf{y} = \mathbf{X}\mathbf{x}$, and thus if our design matrix \mathbf{X} is full-rank then our Hessian is a positive definite matrix by definition.

Answer Q1E: We must show that Newton-Raphson for the two-class logistic regression is equivalent to a form of the IRLS (Iteratively reweighted least squares) algorithm. Newton-Raphson starts with an initial guess of the solution $\hat{\theta}_0$ and recursively updates and iterates until numerical convergence too the solution $\hat{\theta}$. We begin with the l -th iteration:

$$\theta^{(\ell+1)} = \theta^{(\ell)} - (\nabla_{\theta}^2 \log L(\theta|y)|_{\theta=\theta^{(\ell)}})^{-1} \nabla_{\theta} \log L(\theta|y)|_{\theta=\theta^{(\ell)}}. \quad (8)$$

Let us denote $\hat{y}_i \in \mathbb{R}^n$ the vector of conditional probabilities:

$$\hat{\mathbf{y}}_{(l)} = \begin{bmatrix} \sigma(\theta_{(l)}^T \mathbf{x}_1) \\ \vdots \\ \sigma(\theta_{(l)}^T \mathbf{x}_n) \end{bmatrix}$$

Now let $\mathbf{W} \in \mathbb{R}^{n \times n}$ denote the diagonal matrix;

$$\mathbf{W} := \begin{bmatrix} \sigma(\theta_{(l)}^T \mathbf{x}_1)(1 - \sigma(\theta_{(l)}^T \mathbf{x}_1)) & 0 & \cdots & 0 \\ 0 & \sigma(\theta_{(l)}^T \mathbf{x}_2)(1 - \sigma(\theta_{(l)}^T \mathbf{x}_2)) & \cdots & 0 \\ \vdots & \vdots & \ddots & 0 \\ 0 & 0 & \cdots & \sigma(\theta_{(l)}^T \mathbf{x}_n)(1 - \sigma(\theta_{(l)}^T \mathbf{x}_n)) \end{bmatrix}$$

Our design matrix \mathbf{X} is assumed to be a full-rank. Now, the gradient (or score function) and Hessian can be written as follows:

$$\begin{aligned}\nabla_{\theta} \log L(\theta|y) &= \mathbf{X}^T(\mathbf{y} - \mathbf{X}\theta), \\ \nabla_{\theta}^2 \log L(\theta|y) &= -\mathbf{X}^T \mathbf{W} \mathbf{X}.\end{aligned}\tag{9}$$

Substitute this into Equation 8 to give the following:

$$\begin{aligned}\theta^{(\ell+1)} &= \theta^{(\ell)} - (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{X}\theta) \\ &= (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T \mathbf{W} (\mathbf{X}\hat{\theta}_{(\ell-1)} + \mathbf{W}^{-1}(\mathbf{y} - \mathbf{X}\theta)) \\ &= (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T \mathbf{W} \mathbf{z}\end{aligned}\tag{10}$$

Where the adjusted response is given by:

$$\mathbf{z} = \mathbf{X}\theta^{(\ell)} + \mathbf{W}^{-1}(\mathbf{y} - \mathbf{X}\theta).$$

We now can derive the final IRLS equation:

$$\begin{aligned}\theta^{(\ell+1)} &= \underset{\theta}{\operatorname{argmin}} \frac{1}{2} (\mathbf{z} - \mathbf{X}\theta^T)^T \mathbf{W} (\mathbf{z} - \mathbf{X}\theta) \\ &= \underset{\theta}{\operatorname{argmin}} \frac{1}{2} (\tilde{\mathbf{y}}(\theta^{(\ell)}) - \mathbf{X}\theta^T)^T \mathbf{W} (\tilde{\mathbf{y}}(\theta^{(\ell)}) - \mathbf{X}\theta^T), \text{ as required. } \square\end{aligned}\tag{11}$$

PROBLEM 2: CONVEX FUNCTIONS

Definition 2.1 (Convex Function). A function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is convex if its domain is a convex set and for all \mathbf{x}, \mathbf{y} in its domain and all $\lambda \in [0, 1]$ we have:

$$f(\lambda \mathbf{x} + (1 - \lambda)\mathbf{y}) \leq \lambda f(\mathbf{x}) + (1 - \lambda)f(\mathbf{y})$$

By definition 2.1 and assuming $A\mathbf{x} + \mathbf{b} \in \text{dom}(f)$, we have:

$$\begin{aligned} g(\lambda \mathbf{x} + (1 - \lambda)\mathbf{y}) &= f(A(\lambda \mathbf{x} + (1 - \lambda)\mathbf{y}) + \mathbf{b}) \\ &= f(\lambda(A\mathbf{x} + \mathbf{b}) + (1 - \lambda)(A\mathbf{y} + \mathbf{b})) \\ &\leq \lambda f(A\mathbf{x} + \mathbf{b}) + (1 - \lambda)f(A\mathbf{y} + \mathbf{b}) \\ &\equiv \lambda g(\mathbf{x}) + (1 - \lambda)g(\mathbf{y}) \\ \implies g(\lambda \mathbf{x} + (1 - \lambda)\mathbf{y}) &\leq \lambda g(\mathbf{x}) + (1 - \lambda)g(\mathbf{y}) \text{ as required. } \quad \square \end{aligned} \tag{12}$$

PROBLEM 3: CONVEX FUNCTIONS

Answer Q3: We are assuming $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is a convex function and \mathbf{x}_* is a local minimum of f . We must show that this local minimum is a global minimizer.

Let our convex domain be \mathcal{D} . Define \mathbf{y} to be any arbitrary vector in our domain, which implies $\mathbf{y} - \mathbf{x}_*$ is well defined. Select an arbitrarily small scalar $\varepsilon > 0$ such that:

$$f(\mathbf{x}_*) \leq f(\mathbf{x}_* + \varepsilon(\mathbf{y} - \mathbf{x}_*)) \tag{13}$$

Given convexity of f we have:

$$f(\mathbf{x}_* + \varepsilon(\mathbf{y} - \mathbf{x}_*)) = f(\varepsilon\mathbf{y} + (1 - \varepsilon)\mathbf{x}_*) \leq \varepsilon f(\mathbf{y}) + (1 - \varepsilon)f(\mathbf{x}_*) \tag{14}$$

Now combining our equations we have the final solution:

$$f(\mathbf{x}_*) \leq \varepsilon f(\mathbf{y}) + (1 - \varepsilon)f(\mathbf{x}_*) \tag{15}$$

which implies $f(\mathbf{x}_*) \leq f(\mathbf{y})$ for any arbitrary \mathbf{y} or $\mathbf{x} \in \mathcal{D}$, thus proving \mathbf{x}_* is a global minimum for our convex f as required. \square

PROBLEM 4: MAJORIZE-MINIMIZATION ALGORITHMS

Definition 4.1 (Majorizer [1]). Suppose g, h are real valued functions on $\mathcal{X} \in \mathbb{R}^n$. For an optimization problem a surrogate function $h(\boldsymbol{\theta}|\bar{\boldsymbol{\theta}})$ is said to be a majorizer of objective function $g(\boldsymbol{\theta})$ provided

$$\begin{aligned} h(\boldsymbol{\theta}|\bar{\boldsymbol{\theta}}) &\geq g(\boldsymbol{\theta}), \text{ for all } \boldsymbol{\theta} \in \mathcal{X}, \\ h(\bar{\boldsymbol{\theta}}|\bar{\boldsymbol{\theta}}) &= g(\bar{\boldsymbol{\theta}}). \end{aligned}$$

This is almost self-evident by definition and is utilized in cases where we optimize over univariate and separable functions in our Majorize-Minimization (MM) algorithms. We are told $g_i(\mathbf{x}|\bar{\mathbf{x}}) \succ f_i(\mathbf{x})$ where we use \succ to denote majorization. Now consider we have

$$g(\mathbf{x}|\bar{\mathbf{x}}) = \sum_{i=1}^n g_i(\mathbf{x}|\bar{\mathbf{x}}), \quad f(\mathbf{x}) = \sum_{i=1}^n f_i(\mathbf{x}).$$

For each separable function in $\{g_i\}_{i=1}^n, \{f_i\}_{i=1}^n$ we have:

$$g_i(\mathbf{x}|\bar{\boldsymbol{\theta}}) \geq f_i(\mathbf{x}), \text{ for all } \mathbf{x} \in \mathcal{R}^p \text{ and } g_i(\bar{\mathbf{x}}|\bar{\mathbf{x}}) = f_i(\bar{\mathbf{x}}).$$

By definition, the sum of these separable functions will also satisfy both constraints in Definition 4.1 implying

$$g(\mathbf{x}|\bar{\mathbf{x}}) = \sum_{i=1}^n g_i(\mathbf{x}|\bar{\mathbf{x}}) \succ f(\mathbf{x}) = \sum_{i=1}^n f_i(\mathbf{x}), \text{ as required. } \quad \square$$

PROBLEM 5: MAJORIZE-MINIMIZATION ALGORITHMS

Answer Q5A: The Hessian that we derived above of the negative log-likelihood in Equation 6 is bounded above by $\frac{1}{4}$ which is similar to the quadratic bounding principle for MM algorithms applied in the paper Böhning, 1992 [2]. Now apply the Taylor's expansion of a 2nd order to define the quadratic function for our MM algorithm: $\ell_i(\theta) \simeq \ell_T(\theta)$ to give:

$$\ell_T(\theta) = \ell_i(\bar{\theta}) - (y_i - \sigma(\bar{\theta}^T \mathbf{x}_i)) \mathbf{x}_i^T (\theta - \bar{\theta}) + \frac{1}{2} (\theta - \bar{\theta})^T \mathbf{x}_i \mathbf{x}_i^T (\theta - \bar{\theta})$$

which implies $g_i(\theta|\bar{\theta})$ is the majorizer of $\ell_i(\theta)$ at $\bar{\theta}$. The MM algorithm then maximizes this quadratic.

Answer Q5B: After using a Taylor series expansion of the second order consider taking the second derivatives. The truncate Taylor expansion of $g(\theta|\bar{\theta})$ gives

$$g(\theta|\bar{\theta}) \approx g(\theta_0) + \nabla g(\theta_0)(\theta - \theta_0) + \frac{1}{2}(\theta - \theta_0)^2 \nabla^2 g(\theta_0)$$

Setting the derivative to zero gives:

$$\frac{\partial g}{\partial \theta} = f(\theta) = \nabla g(\theta_0) + \nabla^2 g(\theta_0)(\theta - \theta_0) = 0$$

Which results in the following update rule:

$$\theta^{(\ell+1)} = \theta^{(\ell)} - \frac{f(\theta_n)}{\nabla f(\theta_n)} = \theta^{(\ell)} - \frac{\nabla g(\theta)}{\nabla^2 g(\theta_n)}. \quad (16)$$

If we pattern match this equation to our solution in Q1E we have the same update rule and thus our update rule for our separable MM algorithm reduces to the same form:

$$\theta^{(\ell+1)} = \underset{\theta}{\operatorname{argmin}} \frac{1}{2} (\tilde{\mathbf{y}}(\theta^{(\ell)}) - \mathbf{X}\theta^T)^T \mathbf{W} (\tilde{\mathbf{y}}(\theta^{(\ell)}) - \mathbf{X}\theta^T), \text{ as required. } \square \quad (17)$$

Answer Q5C: For our general second order Taylor series of our log-likelihood at $\bar{\theta}$ we have:

$$\ell_T(\theta) = \ell_i(\bar{\theta}) + (\theta - \bar{\theta}) \nabla \ell_i(\bar{\theta}) + \frac{1}{2} (\theta - \bar{\theta})^T \nabla^2 \ell_i(\bar{\theta}) (\theta - \bar{\theta}). \quad (18)$$

We are assuming our function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is upper-bounded Hessian meaning there exists $\mathbf{M} = \mu \mathbf{I}$, such that the following holds:

$$\nabla^2 f(\theta) \leq \mu \mathbf{I} \implies \nabla^2 f(\theta) \leq \mathbf{M}. \quad (19)$$

This implies that $M - \nabla^2 f(\theta)$ is non-negative definite for all θ and M serves as our Hessian which is positive semi-definite. Now we can majorize our function using our quadratic upper bound:

$$\begin{aligned} f_T(\theta) &= f(\bar{\theta}) + (\theta - \bar{\theta})^T \nabla f(\bar{\theta}) + \frac{1}{2}(\theta - \bar{\theta})^T \nabla^2 f(\theta - \bar{\theta}) \\ &\leq f_i(\bar{\theta}) + (\theta - \bar{\theta})^T \nabla f(\bar{\theta}) + \frac{1}{2}(\theta - \bar{\theta})^T (\mu \mathbf{I})(\theta - \bar{\theta}) \end{aligned} \tag{20}$$

Erdogdu, Hosseinzadeh, and Zhang, [3] prove that for convex f which has an upper-bounded Hessian $\nabla^2 f \leq \mu \mathbf{I}$ we have the following identity:

$$\frac{1}{2}(\theta - \bar{\theta})^T (\mu \mathbf{I})(\theta - \bar{\theta}) \leq \frac{\mu}{2} \|\theta - \bar{\theta}\|_2^2.$$

Substitute this into our inequality in Equation 20 to give

$$f_T(\theta) \leq f_i(\bar{\theta}) + (\theta - \bar{\theta})^T \nabla f(\bar{\theta}) + \frac{\mu}{2} \|\theta - \bar{\theta}\|_2^2 \equiv \ell(\theta). \tag{21}$$

Thus we have shown that $\ell(\theta)$ is the majorizer for $f(\theta)$ with equality holding when $\theta = \bar{\theta}$.

PROBLEM 6: CLASSIFICATION

Answer Q6A: Prove unconstrained logistic regression does not have a unique solution:

Answer Q6B: Constrained logistic regression with IRLS algorithm with $\beta^{(0)}$ from problem set 1.
Python code below:

```
def logistic_IRLS(X, y, w=None, max_iter=25):
    """ Implement IRLS algorithm for logistic regression

    References:
        [1] Bishop, Pattern Recognition & Machine Learning (2006)

    :X: Design matrix
    :y: Boolean vector of responses
    :w: Initial weight coefficient vector
    :return: w, nll_sequence
    """
    # Set initial weight matrix
    if w is None:
        w = np.array([0]*X.shape[1], dtype='float64')
    y_bar = np.mean(y)
    w_init = math.log(y_bar/(1-y_bar))
    converged = False
    nll_sequence = np.zeros(max_iter)

    for i in range(max_iter):
        h = X.dot(w)
        p = 1/(1+np.exp(-h))
        p_adj = p
        p_adj[p_adj==1.0] = 0.99999999
        nll = -(1 - y.dot(np.log(1-p_adj))) + y.dot(np.log(p_adj))
        nll_sequence[i] = nll
        # Update & check for convergence
        if i>1:
            if not converged and abs(nll_sequence[-1]-nll_sequence[-2])<.000001:
                converged = True
                converged_k = i+1
            elif not converged and np.isnan(nll_sequence[-1]):
                converged = True
                converged_k = i+1

        if not converged:
            s = p*(1-p)
            S = np.diag(s)
            arb_small = np.ones_like(s, dtype='float64')*.000001
            z = h + np.divide((y-p), s, out=arb_small, where=s!=0)
            Xt = np.transpose(X)
            XtS = Xt.dot(S)
            XtSX = XtS.dot(X)
            inverse_of_XtSX = np.linalg.inv(XtSX)
            inverse_of_XtSX_Xt = inverse_of_XtSX.dot(Xt)
            inverse_of_XtSX_XtS = inverse_of_XtSX_Xt.dot(S)
            w = inverse_of_XtSX_XtS.dot(z)
        else:
            nll_sequence[i] = nll
```



```
nll_sequence = np_ffill(nll_sequence,axis=0)
return w, nll_sequence
```

Answer Q6C: Constrained logistic regression with MM algorithm with $\beta^{(0)}$ from problem set 1. Python code below [4]:

```
import numpy as np
from scipy.optimize import minimize_scalar

def logistic(x):
    """ Compute logistic
    """
    return 1 / (1 + np.exp(-x))

def logistic_loss(w, X, y):
    """ Compute logistic loss
    """
    n = X.shape[0]
    z = np.dot(X, w)
    loss = np.sum(np.log(1 + np.exp(-y * z))) / n
    return loss

def logistic_grad(w, X, y):
    """ Compute score (gradient) for logistic Loss
    """
    n = X.shape[0]
    z = np.dot(X, w)
    g = np.dot(X.T, (-y * logistic(-y * z))) / n
    return g

def logistic_majorize(w, X, y, z, L):
    """ MM (Majorize-Minimization) loistic loss approximation
    for logistic regression
    """
    n = X.shape[0]
    mm_loss = logistic_loss(z, X, y) + np.dot(logistic_grad(z, X, y), w - z) + (L /
    2) * np.sum((w - z)**2)

    return mm_loss

def logistic_MM(X, y, w=None, max_iter=100, tol=1e-4, L=0.1):
    """ Implement IRLS algorithm for logistic regression

    References:
        [1] Bishop, Pattern Recognition & Machine Learning (2006)

    :X: Design matrix
    :y: Boolean vector of responses
    :w: Initial weight coefficient vector
    :return: w, nll_sequence
    """
    # Set initial weight matrix
    if w is None:
        w = np.array([0]* X.shape[1], dtype='float64')
    nll_sequence = np.zeros(max_iter)

    for i in range(max_iter):
        z = w.copy()
```

```

res = minimize_scalar(lambda x: logistic_majorize(w, X, y, z, L*x), bounds=(
                                0, 1), method='bounded')
w = z - (1/L) * logistic_grad(z, X, y) / res.x
nll = logistic_loss(w, X, y)
nll_sequence[i] = nll
if np.linalg.norm(w - z) < tol:
    break
return w, nll_sequence

```

Answer Q6D: Plots of negative-log-likelihood for each iteration for IRLS and MM algorithm:

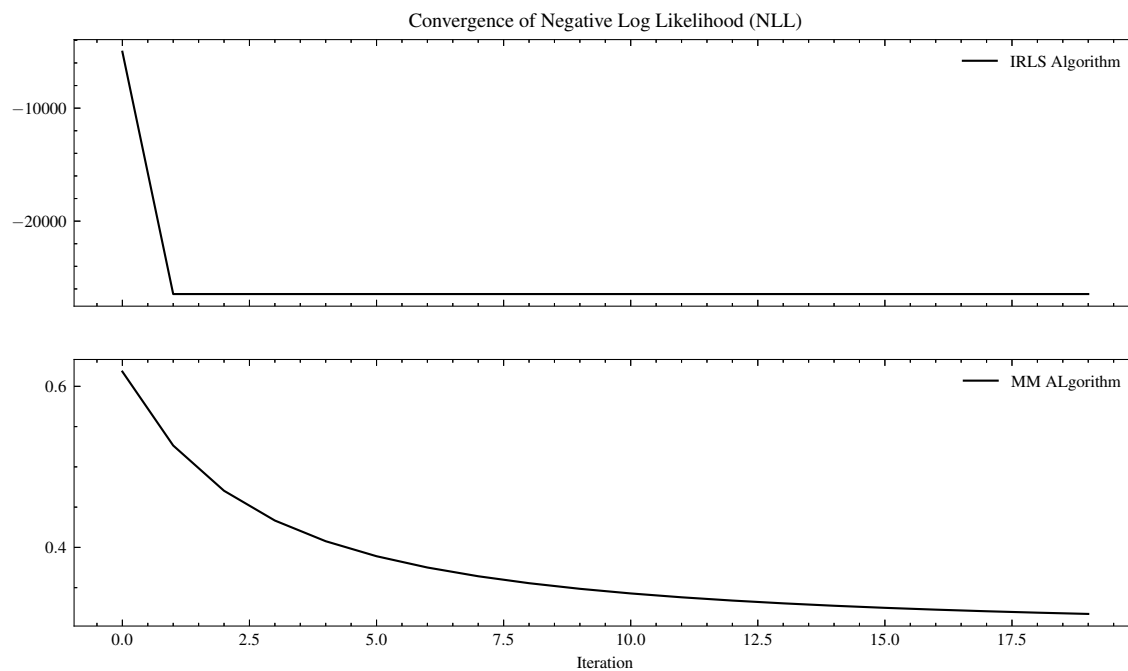


Figure 1: Convergence of NLL for IRLS and MM algorithm

Answer Q6E: Logistic regression solution for IRLS and MM algorithm shown in Fig. 2 below:




Figure 2: Beta Coefficient Plots for IRLS and MM algorithm

Comparison of rankings obtained here with rankings: The IRLS has a similar convergence to our original solution shown in problem set one, it also has a fast convergence. This corresponds with the theoretical convergence shown in 3 steps due to the matrix becoming singular as shown in Bishop and Nasrabadi, 2006 [5]. Our MM algorithm is essentially producing an approximation by utilizing a proxy function and thus the beta coefficient weights are slightly different, but would likely converge after a greater number of iterations. This is similar to ML training that we would see in industry, it is not sufficient just to have a large number of iterations as we have to have a balance between over-fitting our training data due to the Bias-Variance trade-off.

Comment on skill-coefficient vector compared to final NBA team rankings: Our coefficients seem to fairly accurately capture the top-performing (by wins) and worse performing teams (by wins) in this logistic regression framework. Intuitively this makes sense that our parameters accurately reflect the skill of the teams.

PROBLEM 7: CLASSIFICATION

Answer Q7A: Generally, for a multiclass (multinomial) logistic regression with IID data we have:

$$\begin{aligned}
 P(Y = 0|x, \theta) &= \frac{e^{x^T \theta_0}}{1 + \sum_i e^{x^T \theta_i}} \\
 P(Y = 1|x, \theta) &= \frac{e^{x^T \theta_0}}{1 + \sum_i e^{x^T \theta_i}} \\
 &\vdots \\
 P(Y = M-1|x, \theta) &= \frac{e^{x^T \theta_0}}{1 + \sum_i e^{x^T \theta_i}}
 \end{aligned} \tag{22}$$

and thus by the multiplication rule we have:

$$P(Y = M|x, \theta) = \prod_{i=1}^{M-1} \left(\frac{e^{x^T \theta_i}}{1 + \sum_j e^{x^T \theta_j}} \right) \tag{23}$$

Answer Q7B: Not sure - this question is not clear not enough time.

Answer Q7C: Not sure - this question is not clear not enough time.

Answer Q7D: Consider $y = [Y_{i1}, Y_{i2}, \dots, Y_{i(M-1)}]$ is the multinomial trial for dice roll i with $Y_{ij} = 1$ when the response indicates the face of the roll of the dice. Additionally x_i is the vector of explanatory variables and θ is the parameters vector. For simplicity we first consider the log-likelihood for single observation i as follows:

$$\begin{aligned}
 \log L(\theta|(y_i)) &= \log \left[\prod_{m=1}^M \pi_m \right] \\
 &= \sum_{m=1}^{M-1} y_{ij} \log \pi_m + \left(1 - \sum_{m=1}^{M-1} y_{ij}\right) \log \left[1 - \sum_{m=1}^{M-1} \pi_m\right] \\
 &= \sum_{m=1}^{M-1} y_{ij} \log \frac{\pi_m}{1 - \sum_{m=1}^{M-1} \pi_m} + \log \left[1 - \sum_{m=1}^{M-1} \pi_m\right]
 \end{aligned} \tag{24}$$

By definition for our probabilities we have $\pi_M = \sum_{m=1}^{M-1} \pi_m$ and $\sum_{m=1}^M y_{nm} = 1$. Also note that we have N independent observations of dice rolls, then the *negative log-likelihood* (NLL) for the M-class multinomial logistic regression is given as follows:

$$\log L(\theta|(y_{i=1}^n)) = \sum_{i=1}^n \sum_{i=1}^{M-1} -y_{im} + \log \left(1 + \sum_{m'}^{M-1} e^{\theta_{m'}^T x_i} \right) \text{ as required. } \quad \square \tag{25}$$

Answer Q7E: First define the likelihood function as follows:

$$\begin{aligned}
 \log L(\theta | (\mathbf{y}_{i=1}^n)) &= \sum_{m=1}^M y_{nm} \log(\pi_{nm}) \\
 \Rightarrow \nabla_{\theta_m} \log L(\theta | (\mathbf{y}_{i=1}^n)) &= \sum_{m=1}^M y_{nm} \frac{1}{\pi_{nm}} \frac{\partial \pi_{nm}}{\partial t_j} \frac{\partial t_j}{\partial \theta_j} \\
 &= \sum_{m=1}^M y_{nm} y_{nm} \frac{1}{\pi_{nm}} \pi_{nm} (\delta_{mj} - \pi_{nj}) \mathbf{x}_n \\
 &= (y_{nj} - \pi_{nj}) \mathbf{x}_n.
 \end{aligned} \tag{26}$$

As before, given $\sum_{m=1}^M y_{nm} = 1$ then we have:

$$\Rightarrow \nabla_{\theta_m} \log L(\theta | (\mathbf{y}_{i=1}^n)) = - \sum_{i=1}^n x_i (y_{im} - \pi_m), \text{ as required. } \quad \square \tag{27}$$

Answer Q7F: We follow a similar derivation to the highly cited paper by Böhning, 1992 [2] and other texts which used the Kroneker product \otimes to simplify the notation.

Definition 7.1 (Kronecker product). *If A is an $m \times n$ matrix and B is a $p \times q$ matrix, then the Kronecker product $A \otimes B$ is the $pm \times qn$ block matrix:*

$$A \otimes B = \begin{bmatrix} a_{11}B & \dots & a_{1n}B \\ \vdots & \dots & \vdots \\ a_{m1}B & \dots & a_{mn}B \end{bmatrix}$$

Without loss of generality let us reformulate our answer to Q7E as follows for the log-likelihood instead of the negative log-likelihood using Kronecker products:

$$\Rightarrow \nabla_{\theta_m} \log L(\theta | (\mathbf{y}_{i=1}^n)) = \sum_{i=1}^n (\boldsymbol{\pi}_i - \mathbf{y}_i) \otimes \mathbf{x}_i. \tag{28}$$

For $j \in 1, 2, \dots, M-1$ we can also derive the following derivative (the Jacobian of the softmax):

$$\frac{\partial \pi_m}{\partial t_j} = \pi_m (\delta_{mj} - \pi_j) \tag{29}$$

Now our Hessian can be shown to be the following:

$$\begin{aligned}
 \nabla_{\theta}^2 \log L(\theta | (\mathbf{y}_{i=1}^n)) &= \sum_{i=1}^N \mathbf{x}_i \nabla_{\theta_k}^T (\pi_{ij} - y_{ij}) \\
 &= \sum_{i=1}^N \pi_{ij} (\delta_{kj} - \pi_{ik}) \mathbf{x}_i \mathbf{x}_i^T \\
 &\equiv \sum_{i=1}^N (\text{diag}(\boldsymbol{\pi}_i) - \boldsymbol{\pi}_i \boldsymbol{\pi}_i^T) \otimes \mathbf{x}_i \mathbf{x}_i^T \\
 \Rightarrow \nabla_{\theta}^2 \log L(\theta | (\mathbf{y}_{i=1}^n)) &= \sum_{i=1}^N \mathbf{X}_i^T \Sigma_{y_i} \mathbf{X}_i, \text{ as required. } \quad \square
 \end{aligned} \tag{30}$$

REFERENCES

- [1] X.-D. Zhang. “A matrix algebra approach to artificial intelligence”. In: (2020).
- [2] D. Böhning. “Multinomial logistic regression algorithm”. In: *Annals of the institute of Statistical Mathematics* 44.1 (1992), pp. 197–200.
- [3] M. A. Erdogdu, R. Hosseinzadeh, and S. Zhang. “Convergence of Langevin Monte Carlo in chi-squared and Rényi divergence”. In: *International Conference on Artificial Intelligence and Statistics*. PMLR. 2022, pp. 8151–8175.
- [4] G. van den Burg. *Algorithms for Multiclass Classification and Regularized Regression*. Tech. rep. 2018.
- [5] C. M. Bishop and N. M. Nasrabadi. *Pattern recognition and machine learning*. Vol. 4. 4. Springer, 2006.